

Source Data Audit and Final Shona Evaluation Boundary

(a) Strict Unseen Test Construction

Removing training overlap, duplicate identities and label conflicts



Released ShonaSenti test

3,192
rows



Remove training overlap

1,820 rows overlap with released training
under *norm_text* identity

-1,820
rows



Remaining unseen test rows

1,372
rows



Collapse repeated unseen
norm_text identities

1,225
unique identities



Remove label conflict

1 conflicting identity

-1
identity



Strict unseen test

Final evaluation set

1,224
instances

(b) Development Partition Construction

Deduplication, conflict removal and stratified splitting



Released ShonaSenti training

12,767
rows



Normalised-text deduplication

Unique *norm_text* identities

8,781
unique identities



Remove label conflicts

25 conflicting identities

-25
identities



Conflict-free development
pool

8,756
instances



Stratified 90/10 split

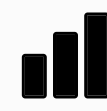
Seed = 42



Build

Candidate discovery,
corpus analysis
and model fitting

7,880
instances



Validation

Threshold,
checkpoint
and hybrid selection

876
instances